

Big data platform final project executive summary

Project overview:

We have developed a kubernetes based application which is designed to store tweets in a database and query this database in order to find meaningful insight into the data.

Data description

We have a CSV file with the tweets of the 20 most popular twitter users. The dataset includes fields such as the tweets author, tweets content, and the data when the tweet was published.

K8 cluster overview and rationale

Our project uses a microservices architecture where each component is encapsulated in its own pod. We have the following 3 separate pods:


- A front end which allows the user to enter the tweet id they wish to perform sentiment analysis on.
- A back end which interacts with our database and performs the sentiment analysis.
- A postgres database which stores our data persistently.

The design rationale behind this configuration is to allow for easy scaling of resources based on demand to a specific service. Deploying each component separately allowed us to work better as a team as each team member could focus on their own portion of work without heavily relying on other parts of code. This is because each microservice can be developed independently of the other.

Key insights from the data

After running sentiment analysis on the tweets we found that many of them had negative sentiment. This confirmed our hypothesis that most people on twitter tweet about negative things rather than positive

Below is a graph showing this for the tweets of Rihanna. As you can see all tweets have a positive sentiment score less than 0.2 out of 1. Clearly she is mostly tweeting about not positive things.



Challenges faced and how we overcame them

While developing and testing the pod which would hold the database everytime we ran the pod it would crash. This presented a problem which at first we struggled to solve. We then used the `kubectl logs` command in order to dive into the problem and understand what was causing the pod to crash. We realized our code which set out the schema for our database was incorrect. After finding this bug by going through the logs we were able to change the lines of code causing it and redeploy the pod with a fixed database image.